

# Explaining neural networks

---

Moritz Wolter

10.03.24

High-Performance Computing and Analytics Lab

# Overview

Linear classifiers

Interpretation by examination

Case Study: Deepfake detection

Input Optimization

The problem with deep CNN

Saliency Maps and Integrated Gradients

Literature

# Motivation

- Neural networks are potent black-box methods.
- Some very deep convolutional neural networks have hundreds of layers and use up to 600mb of disk storage.
- Let's do what we can to open the box!!

## Linear classifiers

---

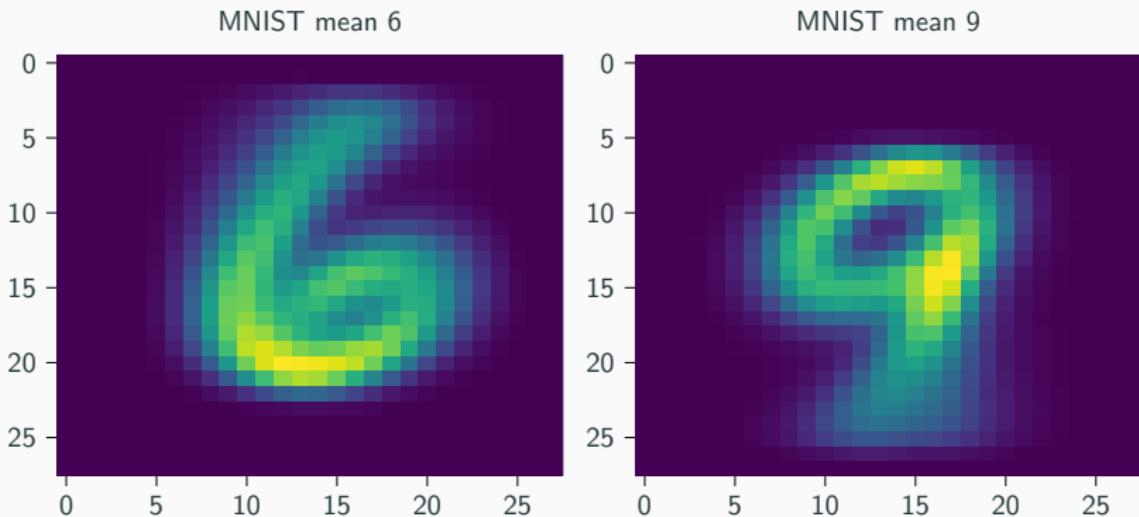
## Definition of a linear classifier

Linear classifiers consist of a dense layer without an activation,

$$\mathbf{o} = \mathbf{Ax} + \mathbf{b}. \quad (1)$$

With  $\mathbf{A} \in \mathbb{R}^{m,n}$ ,  $\mathbf{x} \in \mathbb{R}^n$  and  $\mathbf{b} \in \mathbb{R}^m$ . Linear only works on simple problems that are linearly separable.

# Binary MNIST



## Cross-Entropy

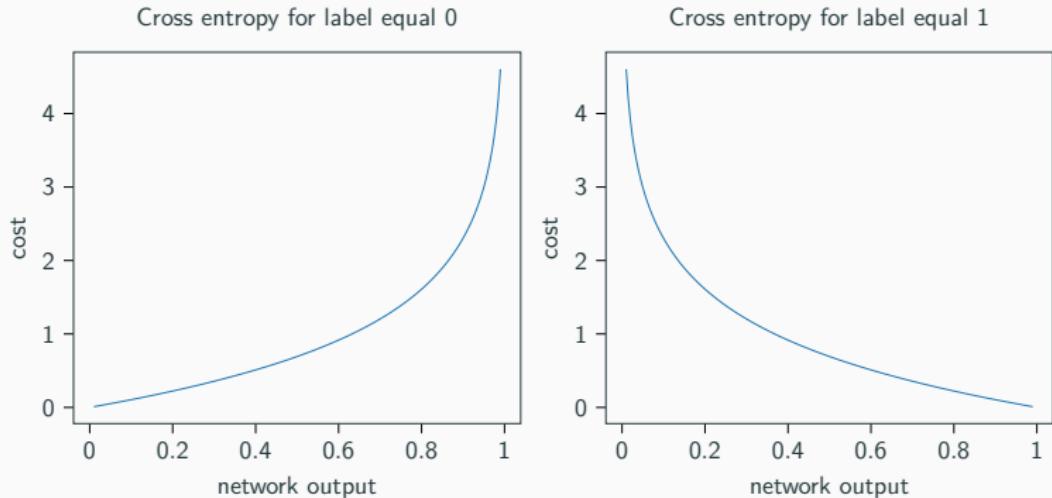
Recall the definition of the cross entropy

$$C_{\text{ce}}(\mathbf{y}, \mathbf{o}) = - \sum_k^{n_o} (\mathbf{y}_k \ln \mathbf{o}_k) + (\mathbf{1} - \mathbf{y}_k) \ln (\mathbf{1} - \mathbf{o}_k). \quad (2)$$

With  $\mathbf{y}, \mathbf{o} \in \mathbb{R}^{n_o}$  defined as vectors of length  $n_o$ .

## Cross-Entropy

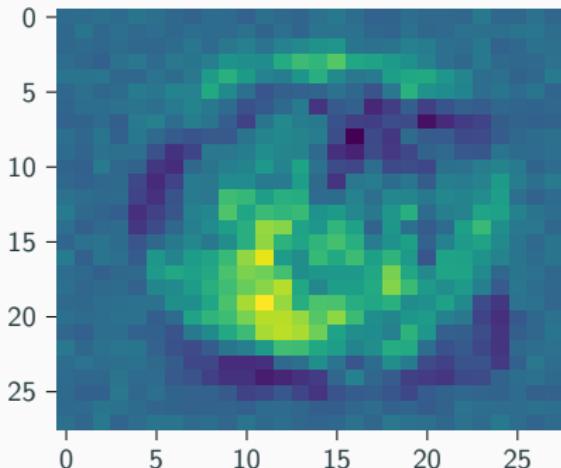
To understand what is going on lets consider the two cases  $y_k = 0$  and  $y_k = 1$ .



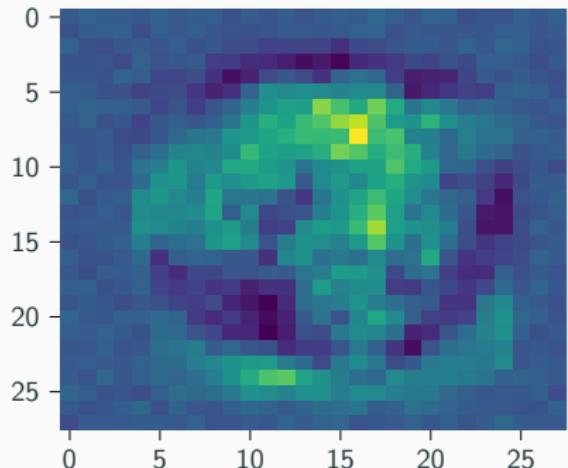
Cross-entropy pushes the output towards the label.

# Interpretation by examination

MNIST binary classifier 6



MNIST binary classifier 9



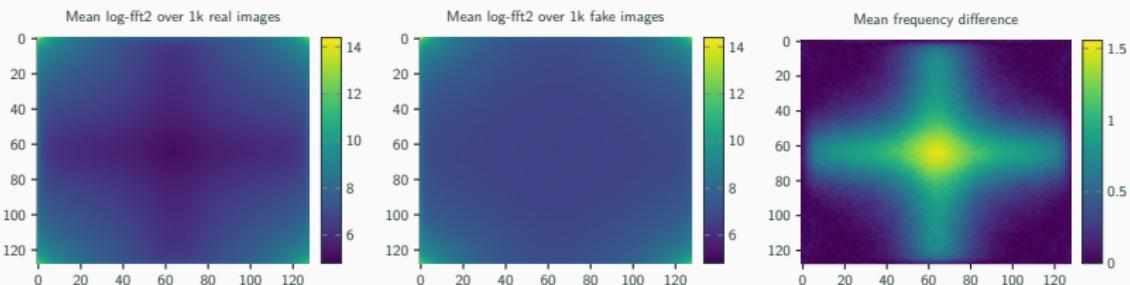
If linear is possible linear is great!!

# What deep-fakes are

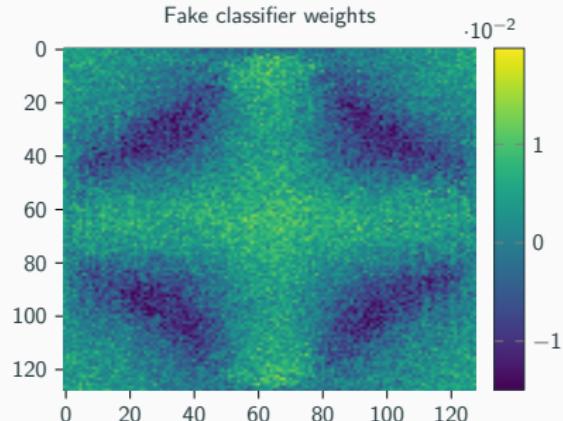
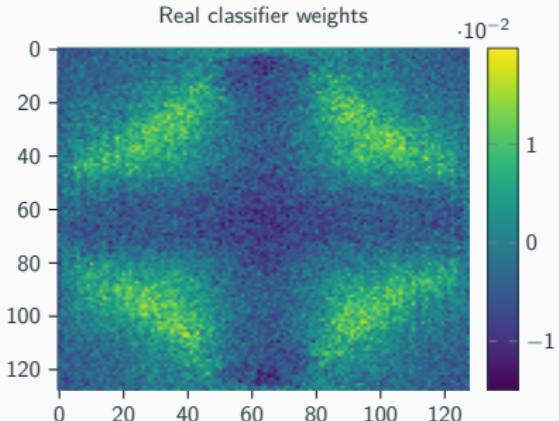
Generative models can generate images. Consider the samples below:



# What blows the con?



# Fake detectors



StyleGAN-generated fakes can be classified with around 99% accuracy this way [Fra+20].

## Summary

- For linearly separable binary problems weight inspection works great!
- Engineered features allow the inspection to reveal something about the data.

## **Input Optimization**

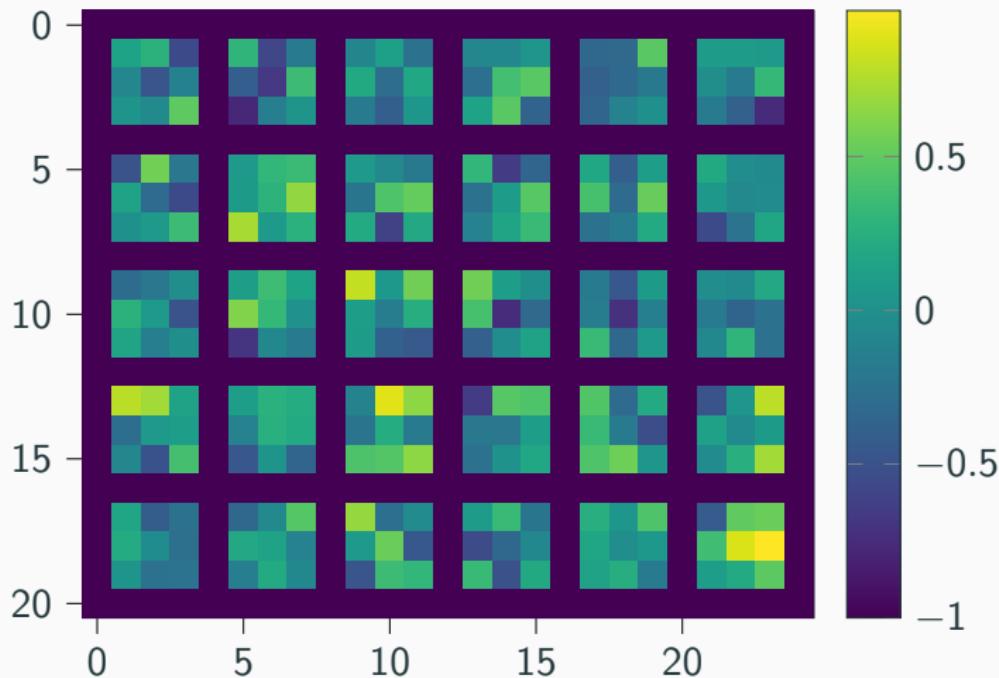
---

# Motivation

Let's be honest. Most linearly separable binary problems are academic.

# An input CNN-Layer

Kernel-shape: (3, 3, 1, 32)



**Figure:** Plot of the input layer kernel weights trained on MNIST.

## Motivation Reloaded

- How do we verify the correct operation of deep nonlinear networks?
- It is *very* hard to interpret the weights of deep networks directly.
- Unit tests would require extensive re-training.

## How to make a single neuron extremely happy

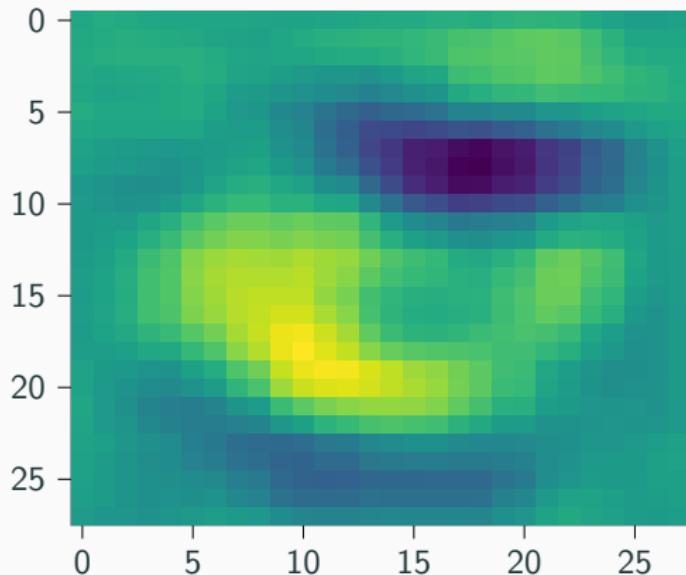
What if we turned the optimization problem around and optimized the input instead of the weights? Consider

$$\max_{\mathbf{x}} y_i = f(\mathbf{x}, \theta), \quad (3)$$

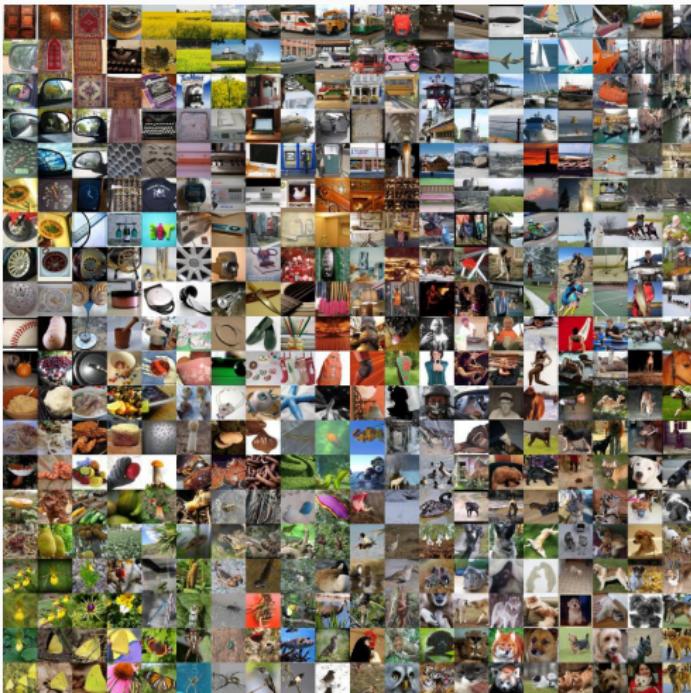
with network weights  $\theta$ , input  $\mathbf{x}$ , and  $y_i$ , the activation of the  $i$ -th output neuron!

## The 6-neurons favorite input

Starting from  $\mathbf{x} = \mathbf{1} \in \mathbb{R}^{1,28,28,1}$  using image  $\mu$ - $\sigma$ -normalization after every step with a step size of one and a positive update yields:



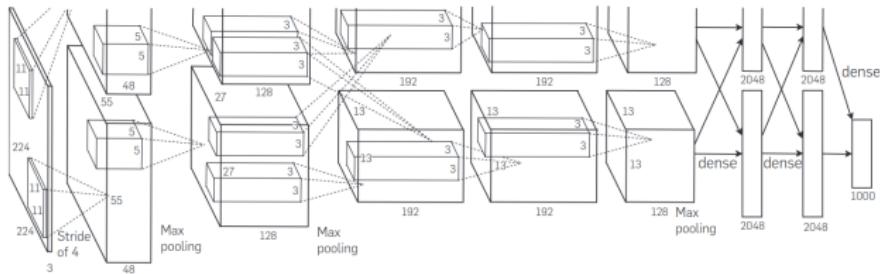
# The image net dataset



**Figure:** Image Net sample images as shown in [Rus+15]. Today 14,197,122 annotated images. Typically with 1000 object categories.

# AlexNet

Figure 2. An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 290,400–186,624–64,896–64,896–43,264–4096–4096–1000.



**Figure:** The Alexnet-architecture used for classify imagenet in 2010  
[KSH17]

## Early Alexnet layers

---

**Figure 3. Ninety-six convolutional kernels of size  $11 \times 11 \times 3$  learned by the first convolutional layer on the  $224 \times 224 \times 3$  input images. The top 48 kernels were learned on GPU 1 while the bottom 48 kernels were learned on GPU 2 (see Section 7.1 for details).**



**Figure:** Plot from [KSH17].

## Saliency Maps

[SVZ13] tell us to optimize

$$\arg \max_{\mathbf{I}} S_c(\mathbf{I}) - \lambda \|\mathbf{I}\|_2^2. \quad (4)$$

With  $S_c$ , the classification score for class c.  $\mathbf{I}$  is the input image, and  $\|\mathbf{I}\|_2$  represents the 2-norm image channels.  $\lambda$  is a regularization parameter.

# CNN-Saliency Map



**Figure:** Input optimization saliency maps of a deep CNN trained on imangenet as shown in [SVZ13].

## Integrated-gradients

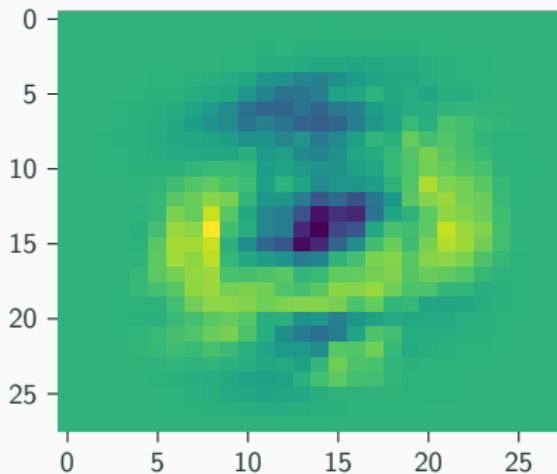
[STY17] propose to estimate individual input contributions to an output neuron via,

$$\text{IntegratedGrads}_i(x) = (x_i - x'_i) \cdot \sum_{k=1}^m \frac{\partial F(x' + \frac{k}{m} \cdot (x - x'))}{\partial x_i}. \quad (5)$$

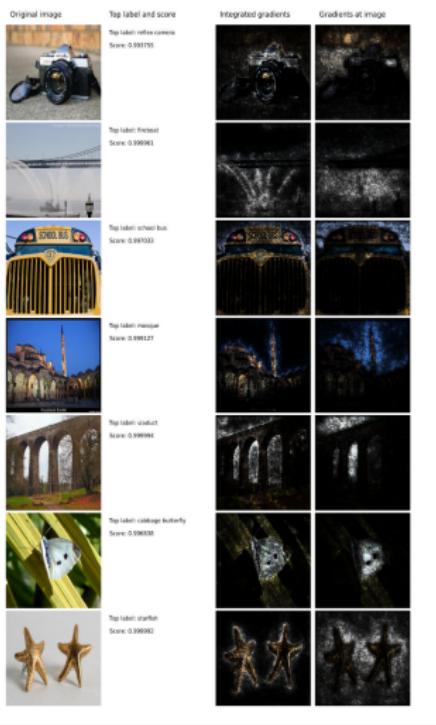
$\frac{\partial F}{\partial x_i}$  denotes the gradients with respect to the input color channels  $i$ .  $x'$  denotes a baseline black image. And  $x$  symbolizes an input we are interested in. Finally,  $m$  denotes the number of summation steps from the black baseline image to the interesting input.

# Integrating the gradients of a multilayer CNN on MNIST

Integrated gradients for the zero neuron on the MNIST-validation set.



# Images from the Wild



**Figure:** Integrated gradient visualization of input saliency for a very deep-CNN trained on Imagenet [Den+09]. Image taken from [STY17].

# Conclusion

- We can look at features and weights and work with input optimization to understand what is going on.

## Literature

---

## References

---

- [Den+09] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. “**Imagenet: A large-scale hierarchical image database.**” In: *2009 IEEE conference on computer vision and pattern recognition*. Ieee. 2009, pp. 248–255.
- [Fra+20] Joel Frank, Thorsten Eisenhofer, Lea Schönherr, Asja Fischer, Dorothea Kolossa, and Thorsten Holz. “**Leveraging frequency analysis for deep fake image recognition.**” In: *International conference on machine learning*. PMLR. 2020, pp. 3247–3258.

- [KSH17] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. “**Imagenet classification with deep convolutional neural networks.**” In: *Communications of the ACM* 60.6 (2017), pp. 84–90.
- [Rus+15] Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, et al. “**Imagenet large scale visual recognition challenge.**” In: *International journal of computer vision* 115 (2015), pp. 211–252.

## Literature iii

- [SVZ13] Karen Simonyan, Andrea Vedaldi, and Andrew Zisserman. “**Deep inside convolutional networks: Visualising image classification models and saliency maps.**” In: *arXiv preprint arXiv:1312.6034* (2013).
- [STY17] Mukund Sundararajan, Ankur Taly, and Qiqi Yan. “**Axiomatic attribution for deep networks.**” In: *International conference on machine learning*. PMLR. 2017, pp. 3319–3328.